Labs
Optimization for Machine Learning
Spring 2019

EPFL
School of Computer and Communication Sciences
Martin Jaggi
github.com/epfml/OptML_course

Problem Set 4, due March 22, 2019
(Proximal Gradient and Subgradient Descent)

Proximal Gradient and Subgradient Descent

Solve Exercises 21, 22, 23, 24 from the lecture notes.

Random Walks

Gradient descent turns up in a surprising number of situations which apriori have nothing to do with optimization. In this exercise, we will see how performing a random walk on a graph can be seen as a special case of gradient descent.

We are given an *undirected* graph $G(V, E)$ with vertices $V = [n]$ labelled 1 through n, and edges $E \subseteq [n]^2$ such that if $(i, j) \in E$, then $(j, i) \in E$. Further, we assume that the graph is *regular* in the sense that every edge has the same degree. Let d be the degree of each node such that if we denote $\mathcal{N}(i) = \{j : (i, j) \in E\}$ to be the neighbors of i, then $|\mathcal{N}(i)| = d$. We assume that every node is connected to itself and so $(i, i) \in \mathcal{N}(i)$.

Now we start our random walk from node 1, jumping randomly from a node to its neighbor. More precisely, suppose at time step t we are at node i_t. Then i_{t+1} is picked uniformly at random from $\mathcal{N}(i)$. If we run this random walk for a large enough T steps, we expect that $\Pr(i_T = j) = 1/n$ for any $j \in [n]$. This is called the stationary distribution.

Problem A. Let us represent the position at time step t in the graph with $\mathbf{e}_{i_t} \in \mathbb{R}^n$ where the i_tth coordinate is 1 and all others are 0. Then, the vector $\mathbf{x}_t = \mathbb{E}[\mathbf{e}_{i_t}]$ denotes the probability distribtion over the n nodes of the graph. Further, let us denote $\mathbf{G} \in \mathbb{R}^{n \times n}$ be the transition probability matrix such that

$$\mathbf{G}_{i,j} = \begin{cases} \frac{1}{d} & \text{if } (i, j) \in E \\ 0 & \text{otherwise}. \end{cases}$$

Show that

$$\mathbf{x}_{t+1} = \mathbf{G}\mathbf{x}_t \tag{1}$$

Problem B. Simulate the random walk above over a torus and confirm that we indeed converge to a uniform distribution over the nodes. What is the *rate* at which this convergence occurs?

Follow the Python notebook provided here:

github.com/epfml/OptML_course/tree/master/labs/ex04/

Problem C. Define $\boldsymbol{\mu} := \frac{1}{n}\mathbf{1}_n$ be a vector of all $1/n$, and a objective function $f : \mathcal{S} \to \mathbb{R}$ as

$$f(\mathbf{x}) = (\mathbf{x} - \boldsymbol{\mu})^\top (\mathbf{I} - \mathbf{G})(\mathbf{x} - \boldsymbol{\mu}),$$

defined over the subspace $\mathcal{S} \subseteq \mathbb{R}^n$ where $\mathcal{S} = \{\mathbf{v} : \mathbf{1}_n^\top \mathbf{v} = 1\}$.

1. Show that f defined above is convex and compute its smoothness constant.

2. Show that running gradient descent on f with the correct step-size is equivalent to the random walk step (1).

3. Prove that $\mathbf{x}_t$ converges to the distribution $\boldsymbol{\mu}$ at a linear rate i.e. for the random walk on a torus with n nodes,

$$\|\mathbf{x}_t - \boldsymbol{\mu}\|_2^2 \leq \left(1 - \frac{1}{n}\right)^t \|\mathbf{x}_0 - \boldsymbol{\mu}\|_2^2 \leq \left(1 - \frac{1}{n}\right)^t.$$

Hint: Use that the two largest eigenvalues of $\mathbf{G}$ are 1 and $1 - \frac{1}{n}$. Also $\mathbf{G}\boldsymbol{\mu} = \boldsymbol{\mu}$ and so $\boldsymbol{\mu}$ is the eigenvector corresponding to eigenvalue 1.