

# Recent Policy Gradient and Actor Critic Methods

We will be looking at four different policy gradient and actor critic methods :

- **TRPO** ('Trust Region Policy Optimization')
- **PPO** ('Proximal Policy Optimization')
- **ACKTR** (Kronecker-factored approximation)
- **PCL** ('Path Consistency Learning')

These methods are all compatible with neural network function approximation.

# Trust Region Policy Optimization - 1 (Schulman et al. 2015)

- Goal = stabilizing (parameterized) policy updates.
- The high-level idea is to take steps in directions that improve the policy, while simultaneously not straying too far from the old policy.
- Making too large a change from the previous policy, especially in high-dimensional, nonlinear environments, can lead to a dramatic decrease in performance. For example, a little forward lean helps running speed, but too much forward lean leads to a crash.
- TRPO takes a principled approach to controlling the rate of policy change - the algorithm places a constraint on the average Kullback-Leibler divergence between the new and old policy after each update.

## Trust Region Policy Optimization - 2

- The change in reward  $\eta$  by updating policy  $\pi$  to policy  $\hat{\pi}$  is

$$\eta(\hat{\pi}) = \eta(\pi) + \sum_s \rho_{\hat{\pi}}(s) \sum_a \hat{\pi}(a|s) A_{\pi}(s, a) \quad (1)$$

- The dependency on  $\hat{\pi}$  through  $\rho_{\hat{\pi}}(s)$  is complex, so we approximate/linearize via

$$L_{\pi}(\hat{\pi}) = \eta(\pi) + \sum_s \rho_{\pi}(s) \sum_a \hat{\pi}(a|s) A_{\pi}(s, a) \quad (2)$$

- In this case, one proves that by taking linear mixture steps, we can guarantee monotonic policy improvement :

$$\pi_{new}(a|s) = (1 - \alpha)\pi_{old}(a|s) + \alpha \arg \max L_{\pi_{old}}(\pi)(a|s)$$

## Theorem (Schulman, 2015):

$$\eta(\hat{\pi}) \geq L_{\pi}(\hat{\pi}) - \frac{4\gamma}{(1-\gamma)^2} \cdot \max_{s,a} |A_{\pi}(s,a)| \cdot D_{KL}^{\max}(\pi, \hat{\pi})$$

(proven using policy coupling methods, and then comparing total variation distance and KL). Based on this pessimistic surrogate, we could iterately solve for policies ( $\pi_i$ ):

$$\pi_{i+1} = \arg \max_{\pi} [L_{\pi_i}(\pi) - C \cdot D_{KL}^{\max}(\pi_i, \pi)]$$

This is a penalized optimization update.

## Trust Region Policy Optimization - 4

- TRPO is a robust approximation to the update on the previous slide, using a **constraint** on the KL divergence rather than a penalty, in order to robustly allow large updates.
- We now use a policy parameter  $\theta$  so that the algorithm becomes:
- Maximize  $L_{\theta_{old}}(\theta)$  over  $\theta$ , subject to constraint  $D_{KL}^{\max}(\theta_{old}, \theta) \leq \delta$ .
- Using Monte-Carlo expectations and cheating a little, this is equivalent to

$$\max_{\theta} \quad \mathbb{E}_{s \sim \rho_{\theta_{old}}, a \sim \theta_{old}} \left[ \frac{\pi_{\theta}(a|s)}{\pi_{\theta_{old}}(a|s)} Q_{\theta_{old}}(s, a) \right]$$

- In practice the algorithm performs the optimization not by SGD, but by **taking a quadratic approximation to the KL-divergence**, and calculating conjugate gradients efficiently (see Hessian-vector products).

- TRPO is useful especially in continuous control tasks, but isn't easily compatible with algorithms that share parameters between a policy and value function or auxiliary losses.
- The algorithm is bespoke as it uses second-order information
- We would like to modify the policy gradients loss function in order to do a **trust region update compatible with stochastic gradient descent**.

## Proximal Policy Optimization - 2

- This happens with the new objective function

$$L^{CLIP}(\theta) = \mathbb{E}_{\text{empirical}} [\min(r_t(\theta), \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)) \cdot A_t]$$

- $r_t$  is the ratio of the probability under the new and old policies, respectively
- $\epsilon$  is a hyperparameter around 0.1 or 0.2.
- The algorithm is simplified : no more KL penalties and adaptive updates. We can train with SGD as usual.
- Why does this work ?

## Proximal Policy Optimization - 3

- For comparison the policy gradient objective is with the same notations

$$L^{PG}(\theta) = \mathbb{E} [\log \pi_{\theta}(a_t|s_t) \cdot A_t]$$

- The TRPO objective is, neglecting the KL term,

$$L^{TRPO}(\theta) = \mathbb{E} \left[ \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)} \cdot A_t \right] = \mathbb{E} [r_t(\theta) \cdot A_t]$$

- In PPO's  $L^{CLIP}$ , the clipping term modifies the surrogate objective by clipping the probability ratio, which removes the incentive for moving  $r_t$  outside of the interval  $[1 - \epsilon, 1 + \epsilon]$ .
- The min term in this scheme is conservative: we only ignore the change in probability ratio when it would make the objective improve, and we include it when it makes the objective worse.



- Idea 1 = policy gradient updates following the **natural gradient**, gives us the direction in parameter space that achieves the largest instantaneous improvement in the objective per unit of change in the output distribution of the network - as measured using the KL-divergence.
- That is, the underlying is the Fisher metric:

$$\mathbb{E}_{\pi} [\nabla_{\theta} \log \pi(a_t|s_t)(\nabla_{\theta} \log \pi(a_t|s_t))^T]$$

- TRPO makes use of Hessian vector-products, in order to avoid inverting the Fisher information matrix explicitly.

- K-FAC (Kronecker-Factored Approximation) is a quasi-Newton algorithm (  $\theta_{t+1} \leftarrow \theta_t - \epsilon H_f^{-1} \nabla_{\theta} f$  ) replacing SGD optimizers for the training of neural networks.
- Idea 2 = use a structured approximation to the inverse Fisher matrix  $H_f^{-1}$ . This makes more progress per step, hence trades sample efficiency for computational cost.
- We can combine both **natural policy gradients** improvement and **K-FAC**.
- Helps solving the sample efficiency problem (faster per step convergence than PG+ADAM).

- Idea = make the (policy) entropic regularization in actor-critic **discounted and recursive** just like the rewards themselves.
- Inserting  $\pi \log \pi$  recursively, this gives the Bellman equation for the loss as

$$L^{ENT}(s, \pi) = \sum_a \pi(a|s) [r(s, a) - \tau \log \pi(a|s) + \gamma L^{ENT}(s', \pi)]$$

- Doing this, all the math works out, with one-hot hard-maxes replaced by Boltzmann soft-maxes (log-sum-exp operator, because it is the **Fenchel-Legendre dual** of the entropy regularization functional).

$$V^*(s) = L^{ENT}(s, \pi^*) = \tau \log \sum_a \exp \left[ \frac{r(s, a) + \gamma V^*(s')}{\tau} \right]$$

$$\pi^*(a|s) = \frac{\exp \left[ \frac{r(s, a) + \gamma V^*(s')}{\tau} \right]}{\exp \left( \frac{V^*(s)}{\tau} \right)}$$

- Hence for each pair  $(s, s')$  and action  $a$  we get at optimality:

$$V^*(s) - \gamma V^*(s') = r(s, a) - \tau \log \pi^*(a|s)$$

- Therefore on any full  $d$ -length sub-trajectory of states  $(s_{i,i+d})$ , we can define the loss  $C(s_{i,i+d}, \theta, \phi)$  as

$$-V_\phi(s_i) + \gamma^d V_\phi(s_{i+d}) + \sum_{j=0}^{d-1} \gamma^j [r(s_{i+j}, a_{i+j}) - \tau \log \pi_\theta(a_{i+j}|s_{i+j})]$$

- The joint  $(\theta, \phi)$  loss

$$L_{\theta, \phi}^{PCL} = \frac{1}{2} \sum_{s_{i,i+d}} C(s_{i,i+d}, \theta, \phi)^2$$

is the MSE objective whose gradients are minimized by PCL.

- Note that under that formulation **policy gradients and (soft) Q-learning are equivalent !**

# DRL is hard... DRL that matters

- Encouraging reproducibility in reinforcement learning
- 'RL algorithms have many moving parts that are hard to debug, and they require substantial effort in tuning in order to get good results.'
- False Discoveries everywhere ? Dependency on the random seed
- Multiple runs are necessary to give an idea of the variability. Reporting the best returns is not enough - at the very least, report top and bottom quantiles
- Don't get discouraged !

Don't be an alchemist

# Suggestions for experiment design

Factors influencing training, by order of importance:

- **Fix your random seed** for reproducibility !
- *Reward scaling* (and more generally reward engineering) help a lot
- *Number of steps* in the calculation of returns
- *Discount factor*, if you have one, also an issue

Clip gradients to avoid NaNs, visualize them with TensorBoard, and finally be **patient** waiting for convergence...

## To go further - References

- Soft Actor Critic
- Boosted Dual Actor Critic
- IMPALA
- etc...